

# Lab 4: Integer Stack Kernel Module Implementation

---

**Amir Gubaidullin CBS-01**

**am.gubaidullin@innopolis.university**

May 4, 2025

## 1. Introduction

Linux kernel module that provides an integer stack data structure accessed through character device. The module supports basic stack operations (push/pop) and stack size configuration via ioctl system calls. User-space utility provides command-line interface for interacting with the kernel module

## 2. Objectives

- Implement a kernel module with dynamic memory allocation
- Implement thread-safe operations using mutex synchronization
- Create a character device driver with file operations
- Develop a user-space CLI utility for device interaction
- Handle edge cases and error conditions properly

## 3. Implementation

### 3.1 I created project structure

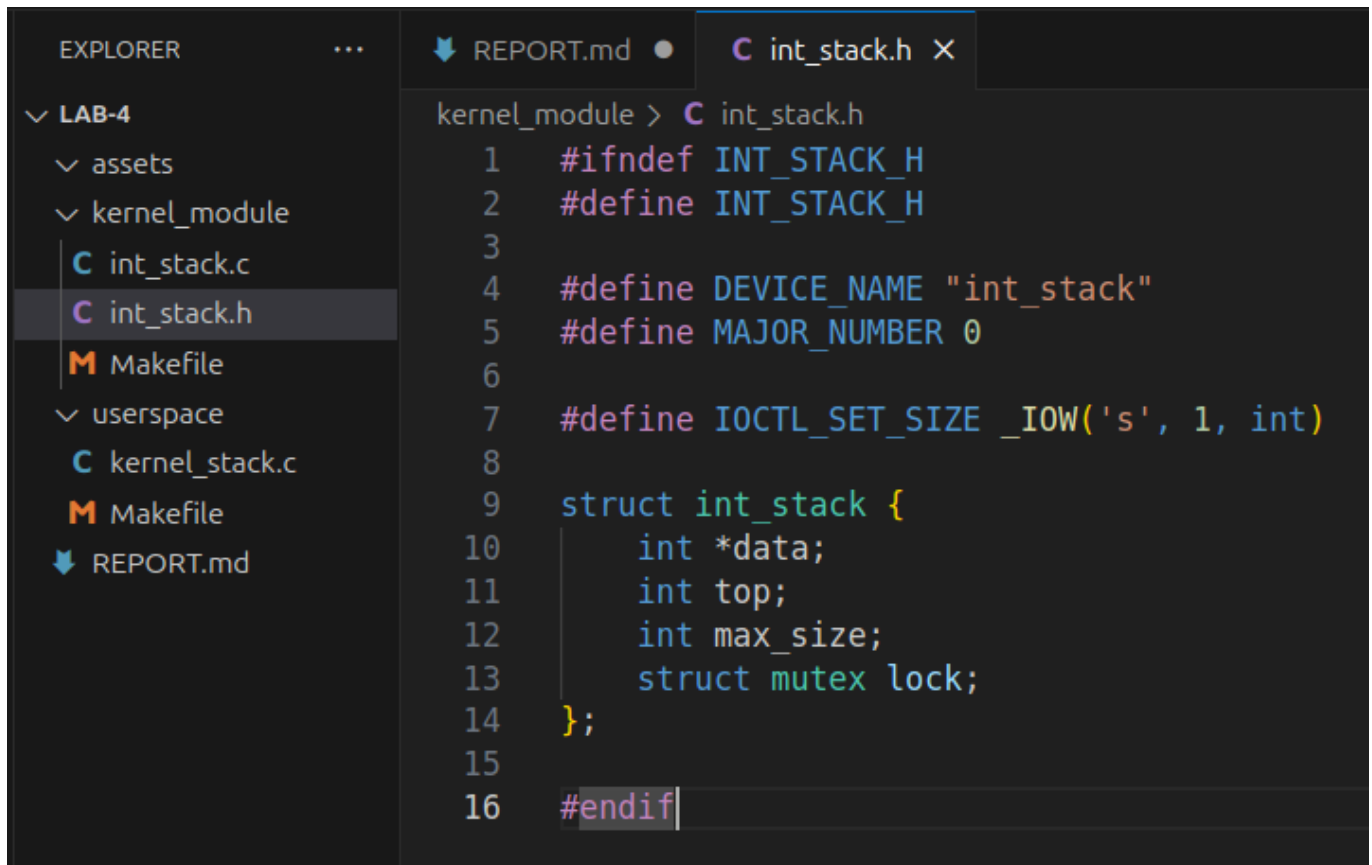
```
mkdir -p kernel_module userspace
```

The project is organized into two main components:

```
lab-4/  
├── kernel_module/  
│   ├── int_stack.h  
│   ├── int_stack.c  
│   └── Makefile  
└── userspace/  
    ├── kernel_stack.c  
    └── Makefile
```

### 3.2 Header file implementation

Header file (int\_stack.h) defines the core data structures, constants:



The screenshot shows a code editor with a dark theme. On the left, the 'EXPLORER' sidebar displays a file tree for 'LAB-4'. The tree includes folders 'assets' and 'kernel\_module', and files 'int\_stack.c', 'int\_stack.h' (selected), 'Makefile', 'userspace', 'kernel\_stack.c', 'Makefile', and 'REPORT.md'. The main editor area shows the content of 'int\_stack.h'. The code defines a device name, major number, and an IOCTL command, and declares a struct for the stack. The struct contains a pointer to an integer array, a top pointer, a maximum size, and a mutex lock.

```
kernel_module > C int_stack.h
1  #ifndef INT_STACK_H
2  #define INT_STACK_H
3
4  #define DEVICE_NAME "int_stack"
5  #define MAJOR_NUMBER 0
6
7  #define IOCTL_SET_SIZE _IOW('s', 1, int)
8
9  struct int_stack {
10     int *data;
11     int top;
12     int max_size;
13     struct mutex lock;
14 };
15
16 #endif
```

Key components:

- Device name definition for registration
- IOCTL command for stack size configuration
- Stack structure with dynamic array, top pointer, maximum size
- Mutex for thread synchronization

### 3.3 Kernel module implementation

The kernel module (int\_stack.c) implements the character device driver:

```

C int_stack.c X
kernel_module > C int_stack.c
1  #include <linux/module.h>
2  #include <linux/kernel.h>
3  #include <linux/init.h>
4  #include <linux/fs.h>
5  #include <linux/device.h>
6  #include <linux/cdev.h>
7  #include <linux/slab.h>
8  #include <linux/uaccess.h>
9  #include <linux/mutex.h>
10
11 #include "int_stack.h"
12
13 MODULE_LICENSE("GPL");
14 MODULE_AUTHOR("Amir Gubaidullin");
15 MODULE_DESCRIPTION("Integer Stack Character Device");
16 MODULE_VERSION("0.1");
17
18 static dev_t dev_num;
19 static struct cdev c_dev;
20 static struct class *cl;
21 static struct int_stack *stack;
22
23 // Function prototypes
24 static int device_open(struct inode *, struct file *);
25 static int device_release(struct inode *, struct file *);
26 static ssize_t device_read(struct file *, char __user *, size_t, loff_t *);
27 static ssize_t device_write(struct file *, const char __user *, size_t, loff_t *);
28 static long device_ioctl(struct file *, unsigned int, unsigned long);
29
30 static struct file_operations fops = {
31     .open = device_open,
32     .release = device_release,
33     .read = device_read,
34     .write = device_write,
35     .unlocked_ioctl = device_ioctl,
36     .owner = THIS_MODULE

```

### 1. Module Initialization (`int_stack_init`):

- Dynamic device number allocation
- Device class creation for automatic device file generation
- Character device registration
- Stack structure allocation and initialization

### 2. File Operations:

- `open()` and `release()`: Basic file operations (currently minimal implementation)
- `read()`: Implements pop operation, returns 0 bytes when stack is empty
- `write()`: Implements push operation, returns -ERANGE when stack is full
- `ioctl()`: Configures stack size, validates input, manages memory allocation

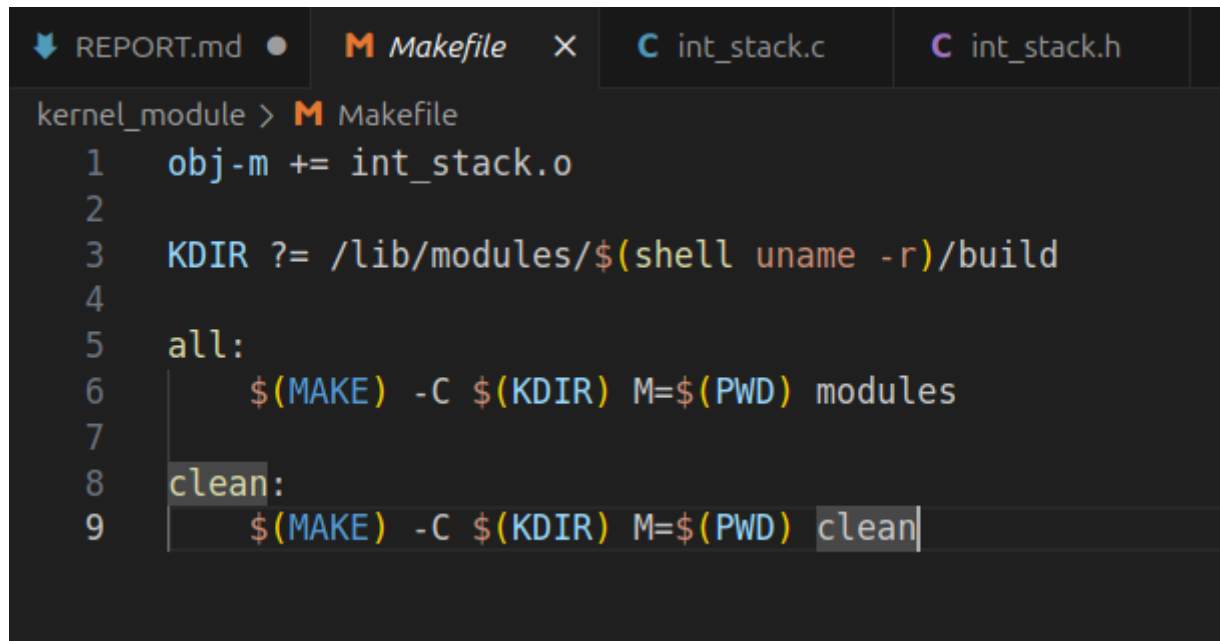
### 3. Thread Safety:

- All stack operations are protected by mutex locks
- Proper lock/unlock sequences in error paths
- Prevents race conditions during concurrent access

### 4. Module Cleanup (`int_stack_exit`):

- Frees allocated memory
- Unregisters character device
- Destroys device class
- Releases device number

Then created kernel module makefile:



The screenshot shows a code editor with four tabs: 'REPORT.md', 'Makefile', 'int\_stack.c', and 'int\_stack.h'. The 'Makefile' tab is active, displaying the following content:

```
kernel_module > M Makefile
1  obj-m += int_stack.o
2
3  KDIR ?= /lib/modules/$(shell uname -r)/build
4
5  all:
6      $(MAKE) -C $(KDIR) M=$(PWD) modules
7
8  clean:
9      $(MAKE) -C $(KDIR) M=$(PWD) clean
```

Features:

- Out-of-tree module compilation
- Automatic kernel version detection
- Clean target for build artifacts

### 3.4 User-space utility

The user-space program (kernel\_stack.c) provides a CLI interface:

```

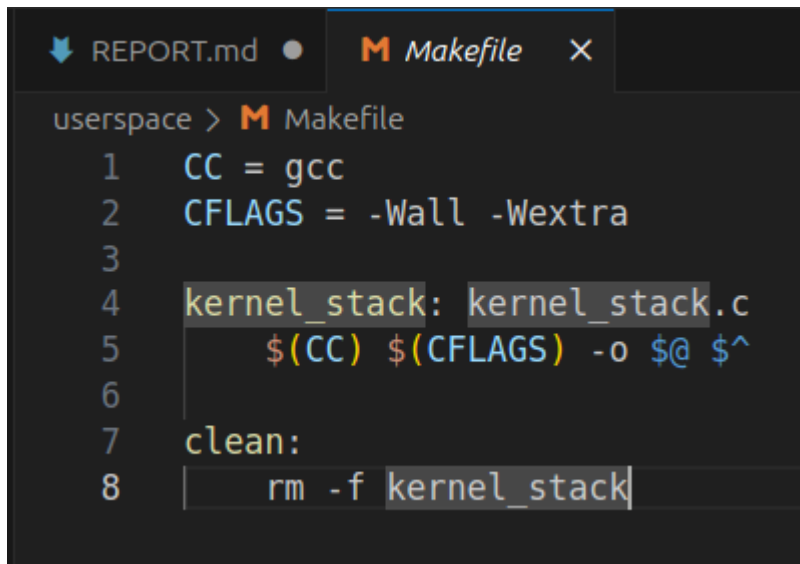
userspace > C kernel_stack.c
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <fcntl.h>
4  #include <unistd.h>
5  #include <sys/ioctl.h>
6  #include <string.h>
7  #include <errno.h>
8
9  #define DEVICE_PATH "/dev/int_stack"
10 #define IOCTL_SET_SIZE _IOW('s', 1, int)
11
12 void print_usage(const char *prog_name) {
13     printf("Usage:\n");
14     printf(" %s set-size <size>\n", prog_name);
15     printf(" %s push <value>\n", prog_name);
16     printf(" %s pop\n", prog_name);
17     printf(" %s unwind\n", prog_name);
18 }
19
20 int main(int argc, char *argv[]) {
21     int fd;
22     int value, size;
23     int ret;
24
25     if (argc < 2) {
26         print_usage(argv[0]);
27         return 1;
28     }
29
30     fd = open(DEVICE_PATH, O_RDWR);
31     if (fd < 0) {
32         perror("Failed to open device");
33         return errno;
34     }
35
36     if (strcmp(argv[1], "set-size") == 0) {
37         if (argc < 3) {

```

Supported commands:

- **set-size <size>**: Configure stack size
- **push <value>**: Add element to stack
- **pop**: Remove and display top element
- **unwind**: Remove and display all elements

Then created user-space makefile:



```

userspace > M Makefile
1  CC = gcc
2  CFLAGS = -Wall -Wextra
3
4  kernel_stack: kernel_stack.c
5      $(CC) $(CFLAGS) -o $@ $^
6
7  clean:
8      rm -f kernel_stack

```

Features:

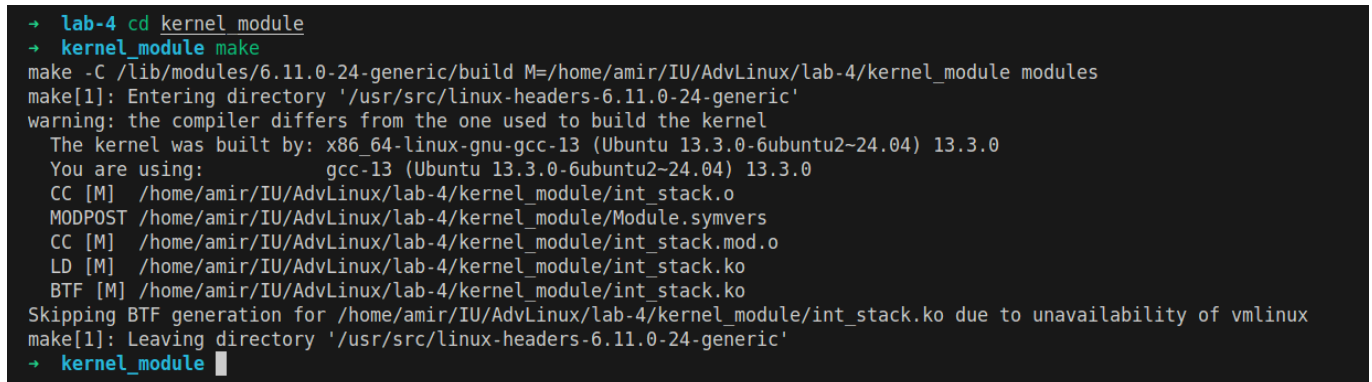
- Strict compilation flags (-Wall -Wextra)
- Simple build process for CLI utility

### 3.5 Building and testing

Successfully built both components:

- Kernel module compiled without warnings
- User-space utility compiled with strict flags

**Built kernel module:**

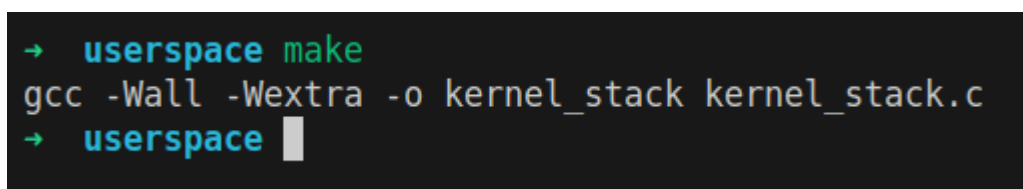


```

→ lab-4 cd kernel_module
→ kernel_module make
make -C /lib/modules/6.11.0-24-generic/build M=/home/amir/IU/AdvLinux/lab-4/kernel_module modules
make[1]: Entering directory '/usr/src/linux-headers-6.11.0-24-generic'
warning: the compiler differs from the one used to build the kernel
The kernel was built by: x86_64-linux-gnu-gcc-13 (Ubuntu 13.3.0-6ubuntu2~24.04) 13.3.0
You are using:          gcc-13 (Ubuntu 13.3.0-6ubuntu2~24.04) 13.3.0
CC [M] /home/amir/IU/AdvLinux/lab-4/kernel_module/int_stack.o
MODPOST /home/amir/IU/AdvLinux/lab-4/kernel_module/Module.symvers
CC [M] /home/amir/IU/AdvLinux/lab-4/kernel_module/int_stack.mod.o
LD [M] /home/amir/IU/AdvLinux/lab-4/kernel_module/int_stack.ko
BTF [M] /home/amir/IU/AdvLinux/lab-4/kernel_module/int_stack.ko
Skipping BTF generation for /home/amir/IU/AdvLinux/lab-4/kernel_module/int_stack.ko due to unavailability of vmlinux
make[1]: Leaving directory '/usr/src/linux-headers-6.11.0-24-generic'
→ kernel_module

```

**Built user-space utility:**



```

→ userspace make
gcc -Wall -Wextra -o kernel_stack kernel_stack.c
→ userspace

```

I faced some issues due to Secure Boot being enabled, the module required signing before loading:

```

# Create signing keys
cd ~

```

```
mkdir kernel-signing
cd kernel-signing
openssl req -new -x509 -newkey rsa:2048 -keyout MOK.priv -outform DER -out
MOK.der -nodes -days 36500 -subj "/CN=Local Kernel Module Signing/"
openssl x509 -in MOK.der -inform DER -outform PEM -out MOK.pem
sudo mokutil --import MOK.der

# Reboot and enroll MOK

# Sign the module
cd /home/amir/IU/AdvLinux/lab-4/kernel_module
sudo /usr/src/linux-headers-$(uname -r)/scripts/sign-file sha256 \
~/kernel-signing/MOK.priv \
~/kernel-signing/MOK.der \
int_stack.ko
```

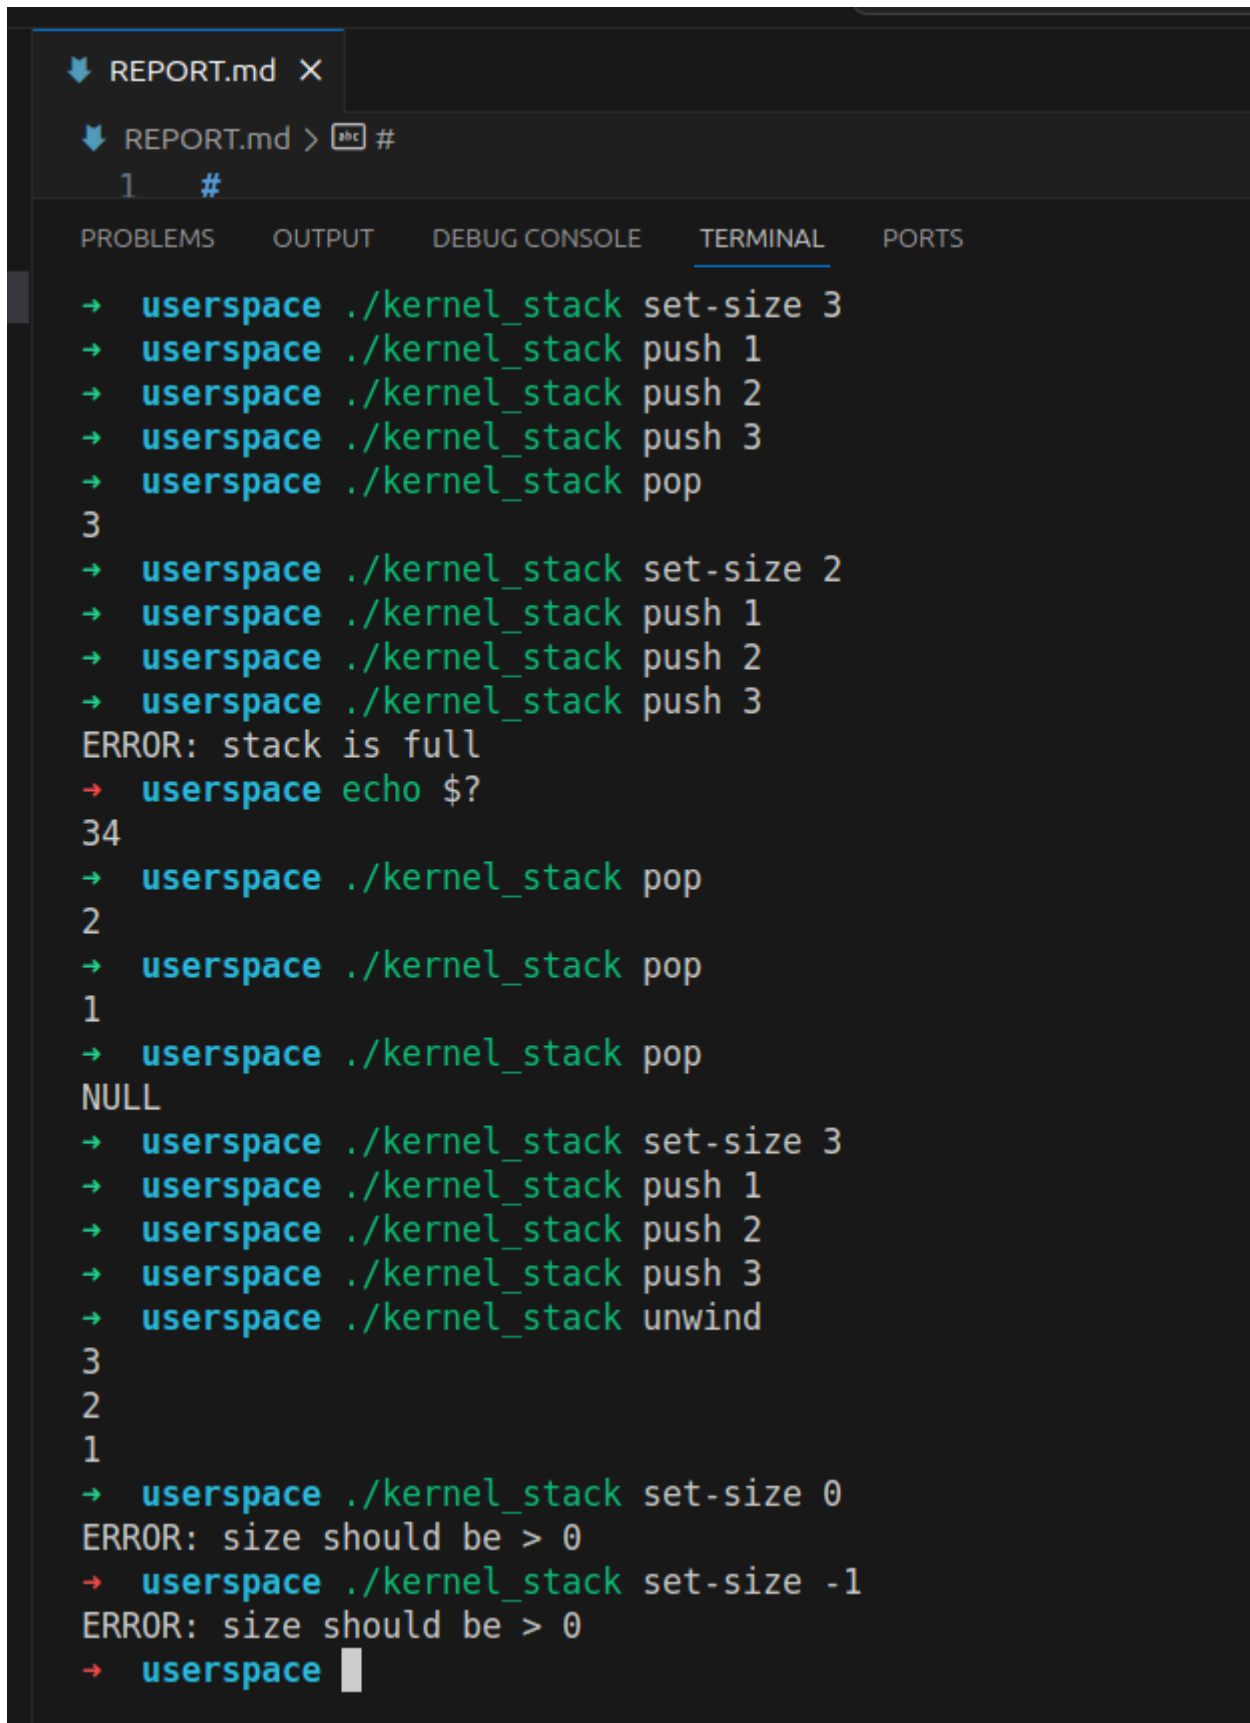
### 3.6 Module loading

After signing, the module was successfully loaded:

```
→ lab-4 cd /home/amir/IU/AdvLinux/lab-4/kernel_module
→ kernel_module sudo /usr/src/linux-headers-$(uname -r)/scripts/sign-file sha256 \
~/kernel-signing/MOK.priv \
~/kernel-signing/MOK.der \
int_stack.ko
→ kernel_module sudo insmod int_stack.ko
→ kernel_module sudo chmod 666 /dev/int_stack
→ kernel_module
```

### 3.7 Test cases

I ran test cases:



```
REPORT.md X
REPORT.md > [abc] #
1 #

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

→ userspace ./kernel_stack set-size 3
→ userspace ./kernel_stack push 1
→ userspace ./kernel_stack push 2
→ userspace ./kernel_stack push 3
→ userspace ./kernel_stack pop
3
→ userspace ./kernel_stack set-size 2
→ userspace ./kernel_stack push 1
→ userspace ./kernel_stack push 2
→ userspace ./kernel_stack push 3
ERROR: stack is full
→ userspace echo $?
34
→ userspace ./kernel_stack pop
2
→ userspace ./kernel_stack pop
1
→ userspace ./kernel_stack pop
NULL
→ userspace ./kernel_stack set-size 3
→ userspace ./kernel_stack push 1
→ userspace ./kernel_stack push 2
→ userspace ./kernel_stack push 3
→ userspace ./kernel_stack unwind
3
2
1
→ userspace ./kernel_stack set-size 0
ERROR: size should be > 0
→ userspace ./kernel_stack set-size -1
ERROR: size should be > 0
→ userspace
```

Executed test cases demonstrating:

### 1. Basic Operations:

- Setting stack size
- Push operations
- Pop operations



- Stack full error handling
- Stack empty handling

## 2. Edge Cases:

- Invalid size values ( $\leq 0$ )
- Stack overflow behavior
- Empty stack pop
- Unwind operation
- Error code verification

## 3.8 Resubmission

Check items are saved when the stack size increases

```

● → kernel_module sudo /usr/src/linux-headers-$(uname -r)/scripts/sign-file sha256 \
  ~/kernel-signing/MOK.priv \
  ~/kernel-signing/MOK.der \
  int_stack.ko
● → kernel_module sudo insmod int_stack.ko
● → kernel_module lsmod | grep int_stack
  int_stack          12288  0
● → kernel_module sudo chmod 666 /dev/int_stack
● → kernel_module cd ../userspace
● → userspace ./kernel_stack set-size 3
● → userspace ./kernel_stack push 1
● → userspace ./kernel_stack push 2
● → userspace ./kernel_stack push 3
● → userspace ./kernel_stack set-size 5
● → userspace ./kernel_stack unwind
  3
  2
  1
○ → userspace █

```

Check for deleting items when reducing the stack size

```

● → userspace ./kernel_stack set-size 5
● → userspace ./kernel_stack push 1
● → userspace ./kernel_stack push 2
● → userspace ./kernel_stack push 3
● → userspace ./kernel_stack push 4
● → userspace ./kernel_stack push 5
● → userspace ./kernel_stack set-size 3
● → userspace ./kernel_stack unwind
3
2
1
○ → userspace █

```

### Check correctness of operations after resizing

```

● → userspace ./kernel_stack set-size 2
● → userspace ./kernel_stack push 10
● → userspace ./kernel_stack push 20
⊕ → userspace ./kernel_stack push 30
ERROR: stack is full
● → userspace ./kernel_stack pop
20
● → userspace ./kernel_stack pop
10
● → userspace ./kernel_stack pop
NULL
○ → userspace █

```

## Appendix

kernel\_module/int\_stack.h: Header file

```

#ifndef INT_STACK_H
#define INT_STACK_H

#define DEVICE_NAME "int_stack"
#define MAJOR_NUMBER 0

#define IOCTL_SET_SIZE _IOW('s', 1, int)

```

```

struct int_stack {
    int *data;
    int top;
    int max_size;
    struct mutex lock;
};

#endif

```

## kernel\_module/int\_stack.c: Kernel module implementation

```

#include <linux/module.h>
#include <linux/kernel.h>
#include <linux/init.h>
#include <linux/fs.h>
#include <linux/device.h>
#include <linux/cdev.h>
#include <linux/slab.h>
#include <linux/uaccess.h>
#include <linux/mutex.h>

#include "int_stack.h"

MODULE_LICENSE("GPL");
MODULE_AUTHOR("Amir Gubaidullin");
MODULE_DESCRIPTION("Integer Stack Character Device");
MODULE_VERSION("0.1");

static dev_t dev_num;
static struct cdev c_dev;
static struct class *cl;
static struct int_stack *stack;

// Function prototypes
static int device_open(struct inode *, struct file *);
static int device_release(struct inode *, struct file *);
static ssize_t device_read(struct file *, char __user *, size_t, loff_t *);
static ssize_t device_write(struct file *, const char __user *, size_t, loff_t *);
static long device_ioctl(struct file *, unsigned int, unsigned long);

static struct file_operations fops = {
    .open = device_open,
    .release = device_release,
    .read = device_read,
    .write = device_write,
    .unlocked_ioctl = device_ioctl,
    .owner = THIS_MODULE,
};

static int device_open(struct inode *inode, struct file *file)

```

```
{
    return 0;
}

static int device_release(struct inode *inode, struct file *file)
{
    return 0;
}

static ssize_t device_read(struct file *filp, char __user *buffer, size_t
length, loff_t *offset)
{
    int value;
    int ret;

    if (length < sizeof(int))
        return -EINVAL;

    mutex_lock(&stack->lock);

    // Pop operation
    if (stack->top == -1) {
        mutex_unlock(&stack->lock);
        return 0; // Return 0 bytes if stack is empty
    }

    value = stack->data[stack->top];
    stack->top--;

    mutex_unlock(&stack->lock);

    ret = copy_to_user(buffer, &value, sizeof(int));
    if (ret)
        return -EFAULT;

    return sizeof(int);
}

static ssize_t device_write(struct file *filp, const char __user *buffer,
size_t length, loff_t *offset)
{
    int value;
    int ret;

    if (length < sizeof(int))
        return -EINVAL;

    ret = copy_from_user(&value, buffer, sizeof(int));
    if (ret)
        return -EFAULT;

    mutex_lock(&stack->lock);

    // Push operation
```

```
if (stack->max_size > 0 && stack->top >= stack->max_size - 1) {
    mutex_unlock(&stack->lock);
    return -ERANGE; // Stack is full
}

stack->top++;
stack->data[stack->top] = value;

mutex_unlock(&stack->lock);

return sizeof(int);
}

static long device_ioctl(struct file *file, unsigned int cmd, unsigned long
arg)
{
    int size;
    int ret;

    switch (cmd) {
    case IOCTL_SET_SIZE:
        ret = copy_from_user(&size, (int __user *)arg, sizeof(int));
        if (ret)
            return -EFAULT;

        if (size <= 0)
            return -EINVAL;

        mutex_lock(&stack->lock);

        // Free old buffer if exists
        if (stack->data) {
            kfree(stack->data);
            stack->data = NULL;
        }

        // Allocate new buffer
        stack->data = kmalloc(size * sizeof(int), GFP_KERNEL);
        if (!stack->data) {
            mutex_unlock(&stack->lock);
            return -ENOMEM;
        }

        stack->max_size = size;
        stack->top = -1; // Reset stack

        mutex_unlock(&stack->lock);

        return 0;

    default:
        return -EINVAL;
    }
}
```

```
static int __init int_stack_init(void)
{
    // Allocate device number
    if (alloc_chrdev_region(&dev_num, 0, 1, DEVICE_NAME) < 0) {
        pr_err("Failed to allocate device number\n");
        return -1;
    }

    // Create device class
    cl = class_create("chardrv");
    if (IS_ERR(cl)) {
        unregister_chrdev_region(dev_num, 1);
        pr_err("Failed to create class\n");
        return PTR_ERR(cl);
    }

    // Create device
    if (device_create(cl, NULL, dev_num, NULL, DEVICE_NAME) == NULL) {
        class_destroy(cl);
        unregister_chrdev_region(dev_num, 1);
        pr_err("Failed to create device\n");
        return -1;
    }

    // Initialize and add cdev
    cdev_init(&c_dev, &fops);
    if (cdev_add(&c_dev, dev_num, 1) == -1) {
        device_destroy(cl, dev_num);
        class_destroy(cl);
        unregister_chrdev_region(dev_num, 1);
        pr_err("Failed to add cdev\n");
        return -1;
    }

    // Allocate memory for stack structure
    stack = kmalloc(sizeof(struct int_stack), GFP_KERNEL);
    if (!stack) {
        cdev_del(&c_dev);
        device_destroy(cl, dev_num);
        class_destroy(cl);
        unregister_chrdev_region(dev_num, 1);
        pr_err("Failed to allocate stack structure\n");
        return -ENOMEM;
    }

    // Initialize stack
    stack->data = NULL;
    stack->top = -1;
    stack->max_size = 0;
    mutex_init(&stack->lock);

    pr_info("int_stack: module loaded\n");
    return 0;
}
```

```

}

static void __exit int_stack_exit(void)
{
    if (stack) {
        if (stack->data)
            kfree(stack->data);
        kfree(stack);
    }

    cdev_del(&c_dev);
    device_destroy(cl, dev_num);
    class_destroy(cl);
    unregister_chrdev_region(dev_num, 1);

    pr_info("int_stack: module unloaded\n");
}

module_init(int_stack_init);
module_exit(int_stack_exit);

```

### kernel\_module/Makefile: Kernel module build configuration

```

obj-m += int_stack.o

KDIR ?= /lib/modules/$(shell uname -r)/build

all:
    $(MAKE) -C $(KDIR) M=$(PWD) modules

clean:
    $(MAKE) -C $(KDIR) M=$(PWD) **clean**

```

### userspace/kernel\_stack.c: User-space utility

```

#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <unistd.h>
#include <sys/ioctl.h>
#include <string.h>
#include <errno.h>

#define DEVICE_PATH "/dev/int_stack"
#define IOCTL_SET_SIZE _IOW('s', 1, int)

void print_usage(const char *prog_name) {
    printf("Usage:\n");
    printf("  %s set-size <size>\n", prog_name);
}

```

```
printf(" %s push <value>\n", prog_name);
printf(" %s pop\n", prog_name);
printf(" %s unwind\n", prog_name);
}

int main(int argc, char *argv[]) {
    int fd;
    int value, size;
    int ret;

    if (argc < 2) {
        print_usage(argv[0]);
        return 1;
    }

    fd = open(DEVICE_PATH, O_RDWR);
    if (fd < 0) {
        perror("Failed to open device");
        return errno;
    }

    if (strcmp(argv[1], "set-size") == 0) {
        if (argc != 3) {
            print_usage(argv[0]);
            close(fd);
            return 1;
        }

        size = atoi(argv[2]);
        if (size <= 0) {
            printf("ERROR: size should be > 0\n");
            close(fd);
            return 1;
        }

        ret = ioctl(fd, IOCTL_SET_SIZE, &size);
        if (ret < 0) {
            perror("ioctl failed");
            close(fd);
            return errno;
        }
    }
    else if (strcmp(argv[1], "push") == 0) {
        if (argc != 3) {
            print_usage(argv[0]);
            close(fd);
            return 1;
        }

        value = atoi(argv[2]);
        ret = write(fd, &value, sizeof(int));
        if (ret < 0) {
            if (errno == ERANGE) {
                printf("ERROR: stack is full\n");
            }
        }
    }
}
```



```

        close(fd);
        return 34; // Return ERANGE
    }
    perror("write failed");
    close(fd);
    return errno;
}
}
else if (strcmp(argv[1], "pop") == 0) {
    ret = read(fd, &value, sizeof(int));
    if (ret == 0) {
        printf("NULL\n");
    } else if (ret < 0) {
        perror("read failed");
        close(fd);
        return errno;
    } else {
        printf("%d\n", value);
    }
}
else if (strcmp(argv[1], "unwind") == 0) {
    while ((ret = read(fd, &value, sizeof(int))) > 0) {
        printf("%d\n", value);
    }
    if (ret < 0) {
        perror("read failed");
        close(fd);
        return errno;
    }
}
else {
    print_usage(argv[0]);
    close(fd);
    return 1;
}

close(fd);
return 0;
}

```

## userspace/Makefile: User-space build configuration

```

CC = gcc
CFLAGS = -Wall -Wextra

kernel_stack: kernel_stack.c
    $(CC) $(CFLAGS) -o $@ $^

clean:
    rm -f kernel_stack

```

